Stochastic optimization
Chance constrained programming

Fabian Bastin
fabian.bastin@umontreal.ca
Université de Montréal – CIRRELT – IVADO – Fin-ML

A long story

- Introduced in 1959 by Charnes and Cooper
  `https://dl.acm.org/doi/10.1287/mnsc.6.1.73`
- And also a bit improbable.
- Cooper dropped high-school to support his family, and became a professional boxer.
- Became an accountant for Eric Louis Kohler, met while hitchhiking.
- Kohler financed his bachelor at University of Chicago.
- At 26, he enrolled at Columbia University and finished his coursework and dissertation, but never received his PhD due to its clam that decision making was not a centralized process.

A long story (cont'd)

- The collaboration with Charnes was however successful, with more than 200 publications, and led a successful academic carrer.
- Source: `https://www.informs.org/Explore/History-of-O.R.-Excellence/Biographical-Profiles/Cooper-William-W`

Cooper and Charnes

INFORMS John Von Neumann prize (with Richard J. Duffin)

Motivation

Source: J. Linderoth `https://homepages.cae.wisc.edu/~linderot/classes/ie495/lecture22.pdf`
We consider the toy problem

$$\begin{aligned}
\min_x \ & x_1 + x_2 \\
\text{s.t. } & \xi_1 x_1 + x_2 \geq 7 \\
& \xi_2 x_1 + x_2 \geq 4 \\
& x_1, x_2 \geq 0,
\end{aligned}$$

where $\xi_1 \sim U(1, 4)$, $\xi_1 \sim U(1/3, 1)$.

Instead of requiring that a constraint holds for all the scenarios, we can require a sufficiently large probability to satisfy a constraint.

Chance constraints

1. Separate chance constraints

$$P[\xi_1 x_1 + x_2 \geq 7] \geq \alpha_1$$
$$P[\xi_2 x_1 + x_2 \geq 4] \geq \alpha_2$$

2. Joint (integrated) chance constraint

$$P[\xi_1 x_1 + x_2 \geq 7 \cap \xi_2 x_1 + x_2 \geq 4] \geq \alpha$$

Example: joint chance constraints

$$P[(\xi_1, \xi_2) = (1, 1)] = 0.1 \tag{1}$$
$$P[(\xi_1, \xi_2) = (2, 5/9)] = 0.4 \tag{2}$$
$$P[(\xi_1, \xi_2) = (3, 7/9)] = 0.4 \tag{3}$$
$$P[(\xi_1, \xi_2) = (4, 1/3)] = 0.1 \tag{4}$$

Assume that $\alpha \in (0.8, 0.9]$, and we have the joint constraint

$$P[\xi_1 x_1 + x_2 \geq 7 \cap \xi_2 x_1 + x_2 \geq 4] \geq \alpha$$

We then have to satisfy constraints (2) and (3) and either (1) or (4).

Example: graph

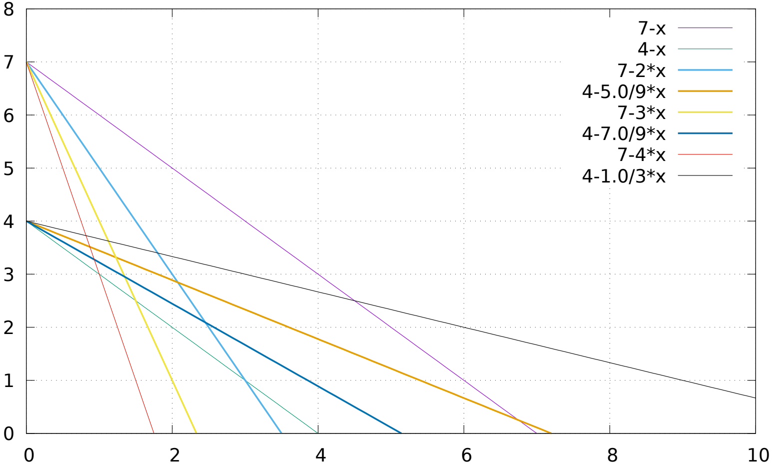

Properties

Feasible set

$$K_1(\alpha) = \{x \mid P[T(\xi)x \geq h(\xi)] \geq \alpha\}$$

$K_1(\alpha)$ is not necessarily convex.

Theorem

Suppose $T(\xi) = T$ is fixed, and $h(\xi)$ has a quasi-concave probability measure P. Then $K_1(\alpha)$ is convex for $0 \leq \alpha \leq 1$.

A function $P : D \to \mathcal{R}$ defined on a domain D is quasi-concave if $\forall$ convex sets $U, V \subseteq D$, and $0 \leq \lambda \leq 1$,

$$P[(1-\lambda)U + \lambda V] \geq \min\{P[U], P[V]\}.$$

Quasi-concave probability distributions

- Uniform

$$f(x) = \begin{cases} 1/\mu(S), & x \in S \\ 0 & \text{otherwise,} \end{cases}$$

  where $\mu(S)$ is the measure of S.

- Exponential density

$$f(x) = \lambda e^{-\lambda x}$$

- Multivariate normal density:

$$f(x) = \frac{1}{\sqrt{(2\pi)^n/2\det(\Sigma)}} e^{-\frac{1}{2}(x-\mu)'\Sigma(x-\mu)}$$

If you have such a density, you can

- use Lagrangian techniques
- use a reduced-gradient technique (see Kall & Wallace, Section 4.1)

Single constraint: easy case

The situation in the single constraint case is somewhat more simple.

Suppose again that $T_i(\xi) = T_i$ is constant. Then

$$P[T_i x \geq h_i(\xi)] = F(T_i x) \geq \alpha$$

so the deterministic equivalent is

$$T_i x \geq F^{-1}(\alpha)$$

...linear constraint! The resulting problem is still linear.

Recall that the inverse of the cdf is defined as

$$F^{-1}(\alpha) = \min\{x : F(x) \geq \alpha\}.$$

Other "solvable" cases

Let $h(\xi) = h$ be fixed, $T(\xi) = (\xi_1, \xi_2, \ldots, \xi_n)$, with $\xi = (\xi_1, \xi_2, \ldots, \xi_n)$ a multivariate normal distribution with mean $\mu = (\mu_1, \mu_2, \ldots, \mu_n)$ and variance-covariance matrix Σ. Then

$$K_1(\alpha) = \{x \mid \mu' x \geq h + \Phi^{-1}(\alpha)\sqrt{x' \Sigma x}\},$$

where Φ is the standard normal cdf.

$K_1(\alpha)$ is a convex set for $\alpha \geq 0.5$.

It is possible to express it as a second order cone constraint:

$$\|\Sigma^{1/2} x\|_2 \leq \frac{1}{\Phi^{-1}(\alpha)}(\mu' x - h)$$

Second-order cone programming

A second-order cone program (SOCP) is a convex optimization problem of the form

$$
\begin{aligned}
\min_{x} \ & f^T x \\
\text{s.t.} \ & \|A_i x + b_i\|_2 \leq c_i^T x + d_i, \ \ i = 1, \ldots, m \\
& Fx = g
\end{aligned}
$$

where $x \in \mathcal{R}^n$, $f, c_i \in \mathcal{R}^n$, $A_i \in \mathcal{R}^{n_i \times n}$, $b_i \in \mathcal{R}^{n_i}$, $d_i \in \mathcal{R}$, $F \in \mathcal{R}^{p \times n}$, and $g \in \mathcal{R}^p$.

SOCPs can be solved by interior point methods.

Example: robust portfolio optimization

(Taken from S. Boyd and J. Linderoth)
Suppose we want to invest in n assets, providing return rates $\beta_1, \beta_2, \ldots, \beta_n$.

The β_i's are random variables. Assume that they are following a multivariate normal distribution with means $\bar{\beta}_i$ and covariance matrix Σ.

Suppose that we want to ensure a return of at least T. We cannot guarantee it all the time, but we want it to occur most of the time.

Example: robust portfolio optimization (cont'd)

Let $x_i \geq 0$ the part of portfolio to invest in stock i. We have the constraints

$$P\left[\sum_{i=1}^{n} \beta_i x_i \geq T\right] \geq \alpha$$

$$\sum_{i=1}^{n} x_i \leq x$$

$$x_i \geq 0,\ i = 1, \ldots, n$$

where x is the total amount to invest.

The chance constraint can be rewritten as

$$\beta' x - \Phi^{-1}(\alpha)\sqrt{x'\Sigma x} \geq T.$$

Example: robust portfolio optimization (cont'd)

We can also interpret x_i as proportion of the portfolio (position of asset i), by normalizing $\|x\|_1$ to 1. T is now the minimum return rate of the portfolio and x is the portfolio allocation.

We can add some constraints on the x_i to ensure diversification. We summarize them by requiring $x \in \mathcal{C}$.

A complete program can now be expressed as

$$\begin{aligned} \max_x \; & E[\beta' x] \\ \text{s.t. } & P\left[\beta' x \geq T\right] \geq \alpha \\ & \sum_{i=1}^{n} x_i = 1 \\ & x \in \mathcal{C} \end{aligned}$$

Example: loss constraint

Setting T to 0 means that we want to ensure that we will no suffer from loss with some probability. Typicially, α is set to 0.9, 0.95, 0.99,...

The chanced-constraint can also be expressed as

$$P\left[\beta' x \leq 0\right] \leq 1 - \alpha = \gamma.$$

We can also allow the sale of some parts of the portfolio by allowing some x_i to be negative.

Numerical illustration

(Taken from S. Boyd – `http://ee364a.stanford.edu/lectures/chance_constr.pdf`)
$n = 10$ assets, $\alpha = 0.95$, $\gamma = 0.05$, $\mathcal{C} = \{x | x \succeq -0.1\}$

Compare

- optimal portfolio
- optimal portfolio without loss risk constraint
- uniform portfolio $(1/n)\mathbf{1}$

portfolio	$E[\beta' x]$	$P[\beta' x \leq 0])$
optimal	7.51	5.0%
w/o loss constraint	10.66	20.3%
uniform	3.41	18.9%

Generalization

A more general form is

$$\begin{aligned} &\min_x \; h(x) \\ &\text{s.t. } P[g_1(x,\xi) \le 0, \ldots, g_r(x,\xi) \le 0] \ge \alpha \\ &\qquad h_1(x) \le 0, \ldots, h_m(x) \le 0. \end{aligned}$$

or

$$\begin{aligned} &\min_x \; h(x) \\ &\text{s.t. } \mathbb{E}\left[\mathcal{I}_{(0,\infty)}\left(g_1(x,\xi) \le 0, \ldots, g_r(x,\xi) \le 0\right)\right] \ge \alpha \\ &\qquad h_1(x) \le 0, \ldots, h_m(x) \le 0, \end{aligned}$$

where

$$\mathcal{I}_{(0,\infty)} = \begin{cases} 1 & \text{if } t \le 0, \\ 0 & \text{if } t > 0, \end{cases}$$

Solution methods for the general case

- Usually very hard.
- Use a bounding approximation or sample average approximation (SAA).
- We will discuss about it in more details when introducing Monte Carlo techniques.

Probabilistic programming

Source: András Prékopa (2003), "Probabilistic Programming", Chapter 5 in "Stochastic Programming", A. Ruszczyński and A. Shapiro (editors), Elsevier.

- Sometimes we only want to maximize a probability.
- General form:

$$\max_{x} P[g_1(x,\xi) \le 0, \ldots, g_r(x,\xi) \le 0]$$
$$\text{subject to } h_1(x) \le 0, \ldots, h_m(x) \le 0.$$

Measures of violation

- A chance constraint also constraint violation with some probability.
- The violation can be large.
- It is often desirable to avoid too large violations.
- Can we penalize the violation?

Value at Risk

Source: `https://web.stanford.edu/class/ee364a/lectures/chance_constr.pdf`

Value-at-risk of random variable Z, at level η:

$$\mathrm{VaR}(Z;\eta) = \inf\{\gamma \mid P[Z \leq \gamma] \geq \eta\}$$

Therefore, the value-at-risk is simply the inverse of the cdf evaluated at η!

$$\mathrm{VaR}(Z;\eta) = F_Z^{-1}(\eta).$$

Conditional Value at Risk

$$\mathrm{CVaR}(Z;\eta) = \inf_{\beta} \left(\beta + \frac{1}{1-\eta} \mathbb{E}\left[(Z-\beta)_+ \right] \right).$$

Assume that the distribution of Z is continuous.

Solution β^* obtained by solving

$$0 = \frac{d}{d\beta} \left(\beta + \frac{1}{1-\eta} \mathbb{E}\left[(Z-\beta)_+ \right] \right) = 1 - \frac{1}{1-\eta} P[Z \geq \beta],$$

leading to

$$P[Z \geq \beta] = 1 - \eta.$$

As Z is continuous, this last equation can be rewritten as

$$P[Z \leq \beta] = \eta = \mathrm{VaR}(Z;\eta).$$

Expected shortfall

Conditional tail expectation (or expected shortfall)

$$\begin{aligned}\mathbb{E}[z \mid z \geq \beta^*] &= \mathbb{E}[\beta^* + (z - \beta^*) \mid z \geq \beta^*] \\ &= \beta^* + \frac{\mathbb{E}\left[(z - \beta^*)_+\right]}{P[z \geq \beta^*]} \\ &= \mathrm{CVaR}(z; \eta)\end{aligned}$$

- Can be added to the objective.
- Can be used as a constraint: *conditional expectation constraint*

$$\mathbb{E}[z \mid z \geq \beta^*] \leq d.$$

Integrated chance constraints

- Consider the stochastic constraints

$$g_i(x, \xi) \leq 0, \ 1, \ldots, r.$$

- Integrated chance constraint:

$$\mathbb{E}\left[\max_i (g_i(x, \xi))_+\right] \leq d.$$

For more details, see Chapter 6, Willem K. Klein Haneveld, Maarten H. van der Vlerk, Ward Romeijnders (2020), "Stochastic Programming - Modeling Decision Problems Under Uncertainty", Springer.